The background of the slide is a dark gray field covered with a complex, interconnected network of thin, light gray lines. At the points where these lines intersect, there are numerous small dots in shades of brown, tan, and black, creating a dense, web-like pattern that resembles a molecular structure or a data network.

# High Performance Computing

## Exercise 1

Annalisa Paladino

# Introduction

## Task:


To estimate the *latency* of various openMPI algorithms for the two collective operations **broadcast** and **scatter** by changing message size, number of MPI processes and allocation type. To build a *performance model* for the latency

## Architecture:

- 2 **THIN** nodes of the **ORFEO** cluster
- 2 sockets per node
- 12 cores per socket




## Methodology:

- **OSU benchmark** to estimate the latency
- Bash scripts to automate data gathering phase
- **SLURM** workload manager
-  for data analysis and models



# Broadcast algorithms

## 1. Basic Linear (bcast 1):



## 2. Chain (bcast 2):






## 3. Split Binary Tree (bcast 4):

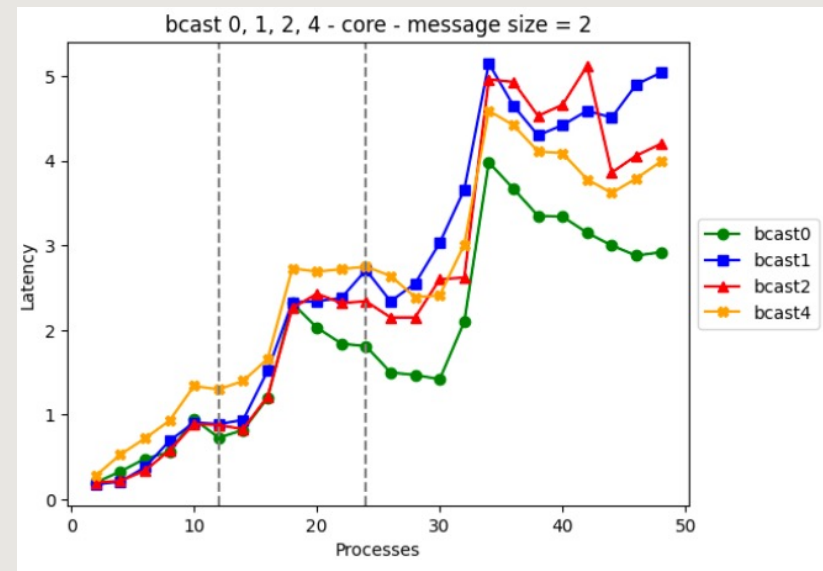
The root process sends the message to two processes, and each internal process thereafter has two children. Before initiating communication, the message is divided in half. The left branch of the tree transmits the left half of the message, and similarly, the right branch transmits the right half. In the final step of the operation, the right and left nodes exchange their respective halves of the message to complete the broadcast

# Bcast: allocation type comparison

## with message size fixed to 2 MPI\_CHAR




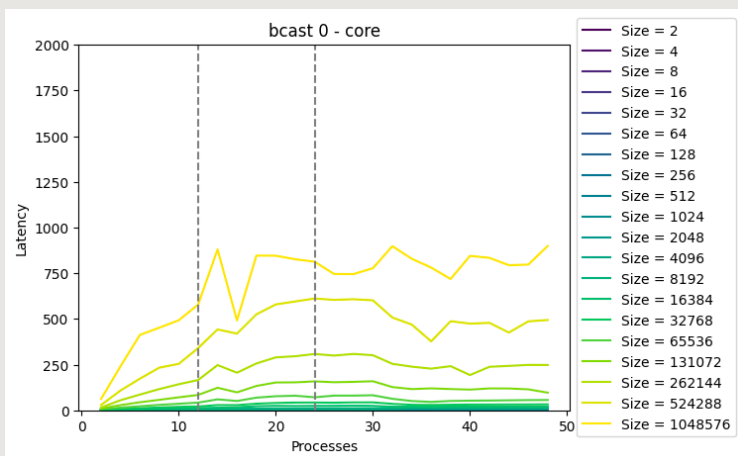


- Socket and core allocations show latency ( $\mu s$ ) “jumps” when changing node, while node allocation progresses seamlessly
- Similar considerations apply to socket changes
- Jumps occur at 16 and 32, hinting at additional factors influencing communication
- Lines convergence: Linear vs Chain algorithm
- Split Binary Tree algorithm (expected) superior performance



# Bcast: algorithms comparison

## non-fixed message size



**Split Binary Tree** algorithm and the default implementation have superior performance:

- Tree structure efficiently distributes the communication among its branches
- Default configurations are robust and efficient in different environments, leveraging platform-specific optimizations

Increasing the message size, **Chain** outperforms **Basic Linear**:

- Chain superior performance, primarily attributable to its utilization of contiguous cores
- Chain capacity to partition messages into chunks during transmission, a feature absent in the linear

# Bcast: performance models

**Initial approach:**  
estimate pt2pt latency  
and bandwidth to  
construct a model  
resembling the  
*Hockney's model*

The model produced for **Basic Linear** (bcast1) is:  
$$\text{latency} = -0.233587 + 0.117043 \times \text{processes}$$
  
with an  $R^2$  adjusted of 0.9381.

The model produced for **Chain** (bcast2) is:  
$$\text{latency} = -0.214384 + 0.109159 \times \text{processes}$$
  
with an  $R^2$  adjusted of 0.8548.




The model produced for **Split Binary Tree** (bcast4) is:  
$$\text{latency} = -0.523 + 5.518 \times \text{processes} - 2.529 \times \text{processes}^2$$
  
with an  $R^2$  adjusted of 0.8793.

# Scatter algorithms


1. **Basic linear** (scatter 1): Each process receives a contiguous block of data from the root process. The root process sends data to each process in sequence.
2. **Binomial** (scatter 2): It uses a binary tree structure to distribute data. Each process receives data from a process that is logarithmically distant from the root process.
3. **Linear non-blocking** (scatter 3): The scatter operations are executed asynchronously. This means that the MPI implementation does not pause for the completion of one scatter operation before initiating the next.

# Scatter: allocation type comparison

## with message size fixed to 2 MPI\_CHAR




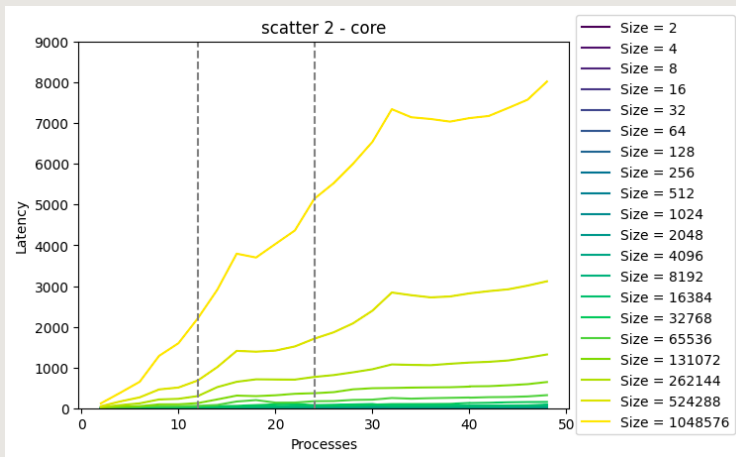
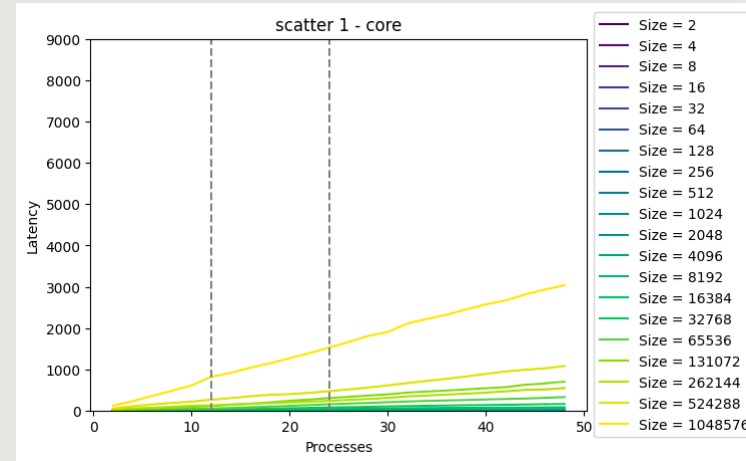
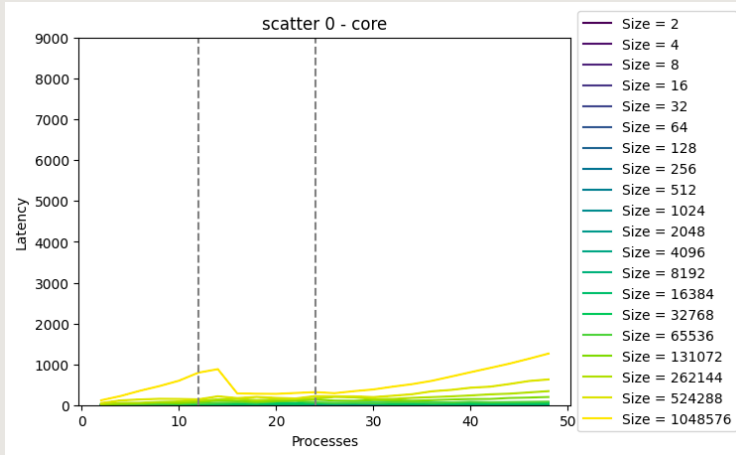
- Jumps occur at 16 and 32, hinting at additional factors influencing communication
- In scatter 2 the latency pattern of node and socket allocations appear closely tied → node-crossing impact socket communication directly
- Scatter 2 shows lower latency values
- In scatter 3 both node and socket show a steady increase until 30 processes and then a subsequent drop





# Scatter: algorithms comparison

## non-fixed message size



**Linear non-blocking** algorithm

outperform the others:

MPI implementation does not pause for the completion of one scatter operation before initiating the next → enable the system to concurrently manage communication and computation tasks

Increasing the message size the **Binomial** algorithm assume the highest values of latency among all the algorithm, while with smaller message sizes it used to be really close to the baseline model

# Scatter: performance models

The model produced for **Basic linear** (Scatter 1) is:  
latency =  $-0.215616 + 0.124241 \times \text{processes}$   
with an  $R^2$  adjusted of 0.9355.

The model produced for **Binomial** (Scatter 2) is:  
latency =  $0.141486 + 0.077624 \times \text{processes}$   
with an  $R^2$  adjusted of 0.7515.

The model produced for **Linear nb** (Scatter 3) is:  
latency =  $-0.278261 + 0.120580 \times \text{processes}$   
with an  $R^2$  adjusted of 0.9317.